

Job Description: Machine Learning Engineer

Company: DataPulse Labs Location: Hybrid (Boston, MA or Remote, EST preferred) Type: Full-

Time Salary Range: \$130,000 - \$170,000 USD annually Posted: May 2025

About Us

DataPulse Labs is a pioneering AI company focused on building intelligent solutions for logistics and supply chain management. We leverage machine learning to optimize processes and deliver actionable insights. Join our innovative team to push the boundaries of AI in real-world applications.

Job Summary

We are seeking a skilled Machine Learning Engineer to develop, train, and deploy advanced ML models for predictive analytics and optimization. You will work with large datasets, design robust algorithms, and ensure models are production-ready. This role focuses on core ML development, requiring deep expertise in data science and model deployment.

Responsibilities

Design and implement machine learning models for use cases like demand forecasting and route optimization using Python and frameworks (TensorFlow, PyTorch, scikit-learn).

Preprocess and clean large datasets to ensure high-quality model inputs.

Train, evaluate, and fine-tune ML models to achieve high accuracy and generalization.

Deploy models to production using MLOps tools (MLflow, Kubeflow) and cloud platforms (AWS, GCP).

Optimize models for performance, scalability, and low latency in real-time applications.

Conduct experiments and A/B testing to validate model performance.

Collaborate with data engineers to build data pipelines for model training and inference.

Monitor and maintain deployed models, addressing drift and retraining as needed.

Document model architectures, experiments, and deployment processes.

Stay current with advancements in ML research and methodologies.

Qualifications

Master's or Ph.D. in Computer Science, Data Science, Statistics, or a related field (or equivalent experience).

3+ years of experience in machine learning model development and deployment.

Expertise in Python and ML libraries (TensorFlow, PyTorch, scikit-learn, pandas).

Strong understanding of supervised and unsupervised learning, including deep learning and ensemble methods.

Experience with cloud platforms (AWS, GCP, or Azure) and containerization (Docker).

Knowledge of MLOps practices for model versioning, monitoring, and automation.

Proficiency in data preprocessing and feature engineering.

Familiarity with big data tools (Spark, Hadoop) is a plus.

Excellent analytical and problem-solving skills.

Ability to communicate technical concepts clearly to diverse stakeholders.

Nice-to-Haves

Experience with time-series forecasting or reinforcement learning.

Familiarity with generative AI or NLP (e.g., transformers, LLMs).

Contributions to ML research or publications in conferences (NeurIPS, ICML).

Knowledge of distributed computing for large-scale model training.

Benefits

Competitive salary with equity options.

Comprehensive medical, dental, and vision coverage.

Hybrid work model with flexible schedules.

Stock options in a fast-growing startup.

20 days PTO + wellness benefits.

Access to AI research resources and professional development funds.

How to Apply

Send your resume, GitHub/Portfolio link, and a cover letter detailing your ML experience to jobs@datapulselabs.com. Highlight any production-grade ML projects or research contributions. DataPulse Labs values diversity and is an equal opportunity employer.